

Efficient AMG preconditioners on clusters of GPUs

Ambra Abdullahi Hassan, Valeria Cardellini, Pasqua D'Ambra, Daniela di Serafino, and
Salvatore Filippone

No Institute Given

Abstract. Many scientific applications require the solution of large and sparse linear systems of equations using iterative methods, this operation accounting for a large percentage of the computing time. In these cases the choice of the preconditioner is crucial for the convergence of the method. Given the broad range of applications, great effort has been put in the development of efficient preconditioners ensuring algorithmic scalability: in this sense multigrid methods have been proved to be particularly promising. Additionally, the advent of GPUs, now found in many of the fastest supercomputers, poses the problem of efficient implementation of these algorithms on highly parallel architectures; this is made more difficult by the fact that the solution of sparse triangular systems, a common kernel in many types of preconditioners, is extremely inefficient on GPUs.

In this paper, we will use the PSBLAS and MLD2P4 libraries to explore the various issues that affect efficiency of multilevel preconditioners on GPUs, both in terms of execution speed and exploitation of computational cores, as well as the algorithmic efficiency in guaranteeing convergence to solution in a number of iterations independent of the degree of parallelism. These issues are investigated in the context of linear systems arising from groundwater modelling application of the filtration of 3D incompressible single-phase flows through anisotropic porous media.

Keywords: