

Efficient Algebraic Multigrid Preconditioners on Clusters of GPUs

Ambra Abdullahi Hassan¹, Valeria Cardellini¹, Pasqua D'Ambra²,
Daniela di Serafino³, and Salvatore Filippone⁴

¹ Università degli Studi di Roma “Tor Vergata”, Roma, Italy
{ambra.abdullahi,cardellini}@uniroma2.it

² Istituto per le Applicazioni del Calcolo “Mauro Picone”, CNR, Napoli, Italy
pasqua.dambra@cnr.it

³ Università degli Studi della Campania “Luigi Vanvitelli”, Caserta, Italy
daniela.diserafino@unicampania.it

⁴ Cranfield University, Cranfield, UK
salvatore.filippone@cranfield.ac.uk

Abstract. Many scientific applications require the solution of large and sparse linear systems of equations using iterative methods, this operation accounting for a large percentage of the computing time. In these cases, the choice of the preconditioner is crucial for the convergence of the iterative method. Given the broad range of applications, great effort has been put in the development of efficient preconditioners ensuring algorithmic scalability: in this sense, multigrid methods have been proved to be particularly promising. Additionally, the advent of GPUs, now found in many of the fastest supercomputers, poses the problem of implementing efficiently these algorithms on highly parallel architectures; this is made more difficult by the fact that the solution of sparse triangular systems, a common kernel in many types of preconditioners, is extremely inefficient on GPUs.

In this paper, we use the PSBLAS and MLD2P4 libraries to explore various issues that affect the efficiency of multilevel preconditioners on GPUs, both in terms of execution speed and exploitation of computational cores, as well as the algorithmic efficiency in guaranteeing convergence to solution in a number of iterations independent of the parallelism degree. We investigate these issues in the context of linear systems arising from groundwater modeling application of the filtration of 3D incompressible single-phase flows through porous media.

1 Introduction

2 Multilevel preconditioner

Multilevel preconditioner e nuclei computazionali dell'algoritmo di applicazione.

3 Sparse linear algebra on GPUs

Sparse linear algebra on GPU e limitazioni del nucleo risoluzione di sistemi triangolari, uso delle inverse approssimate e focus sulla INVK di PSBLAS.

4 Results

Distribuzione dati fatta tenendo conto della geometria (3D) del problema e giustificandola rispetto al vantaggio che fornisce in termini di rapporto calcolo/comunicazione

Acknowledgement

References

- 1.